

# cargo-guppy

track and query Cargo dependencies

Rain <[rain1@fb.com](mailto:rain1@fb.com)>

# The problem

- Analyze Cargo dependency graphs for:
  - license checks
  - dependency audits
  - TCB tracking
  - ...
  - *cool new features*

# Existing solutions

cargo ?

- CLI doesn't have all the features we want
- Rust API isn't for external consumption
  - Large and unstable
  - Many dependencies (e.g. libgit2)
  - Missing documentation

# But...

cargo metadata has:

- package information
- dependency information
- everything we need

```
% cargo metadata --format-version=1 | jq -C | less -FRXK
{
  "packages": [
    {
      "name": "adler",
      "version": "0.2.3",
      "id": "adler 0.2.3 (registry+https://github.com/rust-lang/crates.io-index)",
      "license": "0BSD OR MIT OR Apache-2.0",
      "license_file": null,
      "description": "A simple clean-room implementation of the Adler-32 checksum",
      "source": "registry+https://github.com/rust-lang/crates.io-index",
      "dependencies": [
        {
          "name": "compiler_builtins",
          "source": "registry+https://github.com/rust-lang/crates.io-index",
          "req": "^0.1.2",
          "kind": null,
          "rename": null,
          "optional": true,
          "uses_default_features": true,
          "features": [],
          "target": null,
          "registry": null
        },
        {
          "name": "rustc-std-workspace-core",
          "source": "registry+https://github.com/rust-lang/crates.io-index",
          "req": "^1.0.0",
          "kind": null,
          "rename": "core",
          "optional": true,
          "uses_default_features": true,
          "features": [],
          "target": null,
          "registry": null
        }
      ]
    }
  ]
}
```

The problem

- Dependency audits
- TCB tracking
- Features

solution

cargo ?

- CLI doesn't have all the info
- Rust API isn't for external tools
- Large and unstable
- Many dependencies
- Missing documentation

# Enter guppy

- Read `cargo metadata` input
- Parse as graph structure
- Present nice APIs

# The package graph

- Central structure is `PackageGraph`
- Nodes are packages, edges are dependencies
- Directed, may be cyclic (dev-dependencies)
- Most other types borrow from `PackageGraph`
  - Indicated with a `'g` lifetime
- Uses `petgraph` with integer indexes internally
- Maps integers to borrowed structures externally
- Immutable + `Send` + `Sync` means easy parallelization

# The feature graph

- `FeatureGraph<'g>` is a second, auxiliary graph built from `PackageGraph`
- Nodes are `(package, feature)` pairs, edges are either:
  - Feature dependencies, e.g. `foo = ["bar", "baz"]`
  - Cross-links, e.g. `dep = { version = "1", features = ["foo"] }`
- Computed on-demand

# Core types

| Abstraction            | Package type        | Feature type        |
|------------------------|---------------------|---------------------|
| Main graph             | PackageGraph        | FeatureGraph<'g>    |
| Identifier             | PackageId           | FeatureId<'g>       |
| Extended information   | PackageMetadata<'g> | FeatureMetadata<'g> |
| Dependency edge triple | PackageLink<'g>     | CrossLink<'g> *     |
| Dependency query       | PackageQuery<'g>    | FeatureQuery<'g>    |
| Query result           | PackageSet<'g>      | FeatureSet<'g>      |

\* currently only cross-links are exposed, eventually FeatureLink<'g>



# Core methods

| from           | to            | method               |
|----------------|---------------|----------------------|
| Graph          | Metadata      | metadata             |
| Metadata       | Link iterator | direct_links_        |
| Graph          | Query         | query_               |
| Graph or Query | Set           | resolve_             |
| Set            | Metadata      | packages or features |
| Set            | Link iterator | links                |

\_ indicates that it's several methods, e.g. query\_forward , query\_reverse and query\_directed

# Switching between graphs

| abstraction | $p \sqsubseteq f$ | $f \sqsubseteq p$ |
|-------------|-------------------|-------------------|
| Graph       | feature_graph     | package_graph     |
| Query       | to_feature_query  | to_package_query  |
| Set         | to_feature_set    | to_package_set    |

Package  $\sqsubseteq$  feature requires a `FeatureFilter`. Most people will use `StandardFeatures::None`, `Default` or `All`.

# Filtering during traversals

- Get all transitive dependencies: `PackageQuery::resolve`
- But what if you don't want to follow all edges?
- `PackageQuery::resolve_with()` accepts a `PackageResolver<'g>`
  - Trait with `fn accept(query, link) -> bool`
- Also available as a callback: `PackageQuery::resolve_with_fn`
- Also available for `FeatureQuery`

# Applications

# Basic traversals

- Get all transitive dependencies: `query.resolve()`

- Ignore dev-only dependencies:

```
query.resolve_with_fn(|_, link| !link.dev_only())
```

- Direct dependencies of workspace:

```
query.resolve_with_fn(|_, link| {  
    let (from, to) = link.endpoints();  
    from.in_workspace() && !to.in_workspace()  
})
```

# Cargo builds

- Which packages and features will a build command include?
- Start from a `FeatureSet` describing initials
- Traverse dependency graphs the same way Cargo would
- Customize behavior through `CargoOptions`
  - Platforms and more

```
[metadata]
version = 'v1-install'
include-dev = true
initials-platform = 'proc-macros-on-target'

[metadata.host-platform]
triple = 'thumbv7a-pc-windows-msvc'
target-features = 'all'
flags = ['cargo_web', 'test-flag']

[[metadata.features-only]]
name = 'guppy-benchmarks'
version = '0.1.0'
workspace-path = 'internal-tools/benchmarks'
features = []

[[metadata.features-only]]
name = 'guppy-summaries'
version = '0.2.0'
workspace-path = 'guppy-summaries'
features = []

[[metadata.features-only]]
name = 'proptest-ext'
version = '0.1.0'
workspace-path = 'internal-tools/proptest-ext'
features = []

[[target-package]]
name = 'cargo-compare'
version = '0.1.0'
workspace-path = 'internal-tools/cargo-compare'
status = 'initial'
features = []

[[target-package]]
name = 'fixture-manager'
```

# Cargo builds: v1 and v2 resolvers

- v1 (classic) resolver
  - Packages may or may not be enabled depending on dev, features or platforms
  - Feature resolution is *independent of* which packages are enabled
  - Simulated through 1 feature query + 2 package queries
    - One package query for the target platform, one for the host
- v2 (new) resolver
  - Packages may or may not be enabled depending on dev, features or platforms
  - Feature resolution is *dependent on* which packages are enabled
  - Simulated through 2 feature queries + 2 package queries
    - One each for the target, one each for the host

# Cargo builds: property testing

- Comparison testing with Cargo
  - Generate random queries and compare against Cargo
- Consistency testing with previous versions of guppy
  - Generate random queries and simulate builds
  - Summaries with build results checked into the repo
  - These should only change if there's a good reason



**Cool new features**

# Determinator

- Only run tests for packages that changed from upstream
- Given old metadata, new metadata and paths changed:
  - Map each path to a package
  - Simulate Cargo builds for each package and see which changed
- Support for custom rules
- Diem CI: p25 90% faster, p50 60+%
- [docs.rs/determinator](https://docs.rs/determinator)

```
# Standard ignore and other metadata files.
[[path-rule]]
globs = ["**/.gitignore", "**/.gitattributes", ".dockerignore", ".hgignore", ".svnignore", "**/.ignore"]
mark-changed = []

# Files that can affect the global build. Cargo.toml may contain updates to build flags or profile overrides,
# so rebuild everything if it changes. (We could do a more sophisticated analysis in the future.)
[[path-rule]]
globs = ["rust-toolchain", "Cargo.toml", "**/.cargo/config", "**/.cargo/config.toml"]
mark-changed = "all"

# Tool files that don't influence builds or tests.
[[path-rule]]
globs = ["clippy.toml", "rustfmt.toml", ".lintrules/**/*"]
mark-changed = []

# Cargo.lock is ignored, since the determinator does a deeper analysis to figure out which packages changed.
[[path-rule]]
globs = ["Cargo.lock"]
mark-changed = []

# README, LICENSE and other metadata files are ignored throughout the codebase.
[[path-rule]]
globs = ["**/README*", "**/LICENSE*", "**/CONTRIBUTING*", "**/CODE_OF_CONDUCT*", "**/SECURITY*"]
mark-changed = []
```

# Hakari

- Manage packages for dependency feature unification
  - Workspace-hack packages used by many large projects (rustc, Firefox, Diem)
- Simulate Cargo builds and look for non-workspace packages built with more than one feature set
- Speeds up Diem builds by 15-30% or more
- [docs.rs/hakari](https://docs.rs/hakari)

```
### BEGIN HAKARI SECTION
# version = 'v1'
# verify-mode = true
# unify-target-host = 'none'
# unify-all = false
# platforms = []
#
# [[omitted-packages]]
# name = 'either'
# version = '1.6.1'
# crates-io = true
#
# [[omitted-packages]]
# name = 'petgraph'
# version = '0.5.1'
# crates-io = true
#
# [[omitted-packages]]
# name = 'url'
# version = '2.1.1'
# crates-io = true
#
# [[omitted-packages]]
# name = 'web-sys'
# version = '0.3.45'
# crates-io = true

[target.'cfg(all())'.dependencies]
bstr = { version = "0.2", features = ["default", "lazy_static", "regex-automata", "serde", "serde1", "serde"] }
byteorder = { version = "1", default-features = false, features = ["std"] }
clap = { version = "2", features = ["ansi_term", "atty", "color", "default", "strsim", "suggestions", "vec"] }
itoa = { version = "0.4", features = ["default", "std"] }
libc = { version = "0.2", features = ["default", "std"] }
log = { version = "0.4", default-features = false, features = ["std"] }
memchr = { version = "2", features = ["default", "std", "use_std"] }
num-traits = { version = "0.2", features = ["default", "std"] }
regex = { version = "1", features = ["aho-corasick", "default", "memchr", "perf", "perf-cache", "perf-dfa"] }
regex-syntax = { version = "0.6", features = ["default", "unicode", "unicode-age", "unicode-bol", "unicode-emoji"] }
serde = { version = "1", features = ["default", "derive", "serde_derive", "std"] }
serde_json = { version = "1", features = ["default", "raw_value", "std"] }
winapi = { version = "0.3", default-features = false, features = ["basetsd", "consoleapi", "errhandlingapi"] }

[target.'cfg(all())'.build-dependencies]
libc = { version = "0.2", features = ["default", "std"] }
log = { version = "0.4", default-features = false, features = [] }
proc-macro2 = { version = "1", features = ["default", "proc-macro"] }
quote = { version = "1", features = ["default", "proc-macro"] }
syn = { version = "1", features = ["clone-impls", "default", "derive", "full", "parsing", "printing", "proc-macro"] }

### END HAKARI SECTION
```

Questions?