

cargo-guppy

track and query Cargo dependencies

Rain <rain1@fb.com>

The problem

- Analyze Cargo dependency graphs for:
 - license checks
 - dependency audits
 - TCB tracking
 - ...
 - *cool new features*

Existing solutions

cargo ?

- CLI doesn't have all the features we want
- Rust API isn't for external consumption
 - Large and unstable
 - Many dependencies (e.g. libgit2)
 - Missing documentation

But...

cargo metadata has:

- package information
- dependency information
- everything we need

Enter guppy

- Read `cargo metadata` input
- Parse as graph structure
- Present nice APIs

The package graph

- Central structure is `PackageGraph`
- Nodes are packages, edges are dependencies
- Directed, may be cyclic (dev-dependencies)
- Most other types borrow from `PackageGraph`
 - Indicated with a `'g` lifetime
- Uses `petgraph` with integer indexes internally
- Maps integers to borrowed structures externally
- Immutable + `Send` + `Sync` means easy parallelization

The feature graph

- `FeatureGraph<'g>` is a second, auxiliary graph built from `PackageGraph`
- Nodes are `(package, feature)` pairs, edges are either:
 - Feature dependencies, e.g. `foo = ["bar", "baz"]`
 - Cross-links, e.g. ``dep = { version = "1", features = ["foo"] }`
- Computed on-demand

Core types

Abstraction	Package type	Feature type
Main graph	PackageGraph	FeatureGraph<'g>
Identifier	PackageId	FeatureId<'g>
Extended information	PackageMetadata<'g>	FeatureMetadata<'g>
Dependency edge triple	PackageLink<'g>	CrossLink<'g> *
Dependency query	PackageQuery<'g>	FeatureQuery<'g>
Query result	PackageSet<'g>	FeatureSet<'g>

* currently only cross-links are exposed, eventually FeatureLink<'g>

Core methods

from	to	method
Graph	Metadata	metadata
Metadata	Link iterator	direct_links_
Graph	Query	query_
Graph or Query	Set	resolve_
Set	Metadata	packages or features
Set	Link iterator	links

_ indicates that it's several methods, e.g. query_forward , query_reverse and query_directed

Switching between graphs

abstraction	$p \sqsubseteq f$	$f \sqsubseteq p$
Graph	feature_graph	package_graph
Query	to_feature_query	to_package_query
Set	to_feature_set	to_package_set

Package \sqsubseteq feature requires a `FeatureFilter`. Most people will use `StandardFeatures::None`, `Default` or `All`.

Filtering during traversals

- Get all transitive dependencies: `PackageQuery::resolve`
- But what if you don't want to follow all edges?
- `PackageQuery::resolve_with()` accepts a `PackageResolver<'g>`
 - Trait with `fn accept(query, link) -> bool`
- Also available as a callback: `PackageQuery::resolve_with_fn`
- Also available for `FeatureQuery`

Applications

Basic traversals

- Get all transitive dependencies: `query.resolve()`

- Ignore dev-only dependencies:

```
query.resolve_with_fn(|_, link| !link.dev_only())
```

- Direct dependencies of workspace:

```
query.resolve_with_fn(|_, link| {  
    let (from, to) = link.endpoints();  
    from.in_workspace() && !to.in_workspace()  
})
```

Cargo builds

- Which packages and features will a build command include?
- Start from a `FeatureSet` describing initials
- Traverse dependency graphs the same way Cargo would
- Customize behavior through `CargoOptions`
 - Platforms and more

Cargo builds: v1 and v2 resolvers

- v1 (classic) resolver
 - Packages may or may not be enabled depending on dev, features or platforms
 - Feature resolution is *independent of* which packages are enabled
 - Simulated through 1 feature query + 2 package queries
 - One package query for the target platform, one for the host
- v2 (new) resolver
 - Packages may or may not be enabled depending on dev, features or platforms
 - Feature resolution is *dependent on* which packages are enabled
 - Simulated through 2 feature queries + 2 package queries
 - One each for the target, one each for the host

Cargo builds: property testing

- Comparison testing with Cargo
 - Generate random queries and compare against Cargo
- Consistency testing with previous versions of guppy
 - Generate random queries and simulate builds
 - Summaries with build results checked into the repo
 - These should only change if there's a good reason

Cool new features

Determinator

- Only run tests for packages that changed from upstream
- Given old metadata, new metadata and paths changed:
 - Map each path to a package
 - Simulate Cargo builds for each package and see which changed
- Support for custom rules
- Diem CI: p25 90% faster, p50 60+%
- docs.rs/determinator

Hakari

- Manage packages for dependency feature unification
 - Workspace-hack packages used by many large projects (rustc, Firefox, Diem)
- Simulate Cargo builds and look for non-workspace packages built with more than one feature set
- Speeds up Diem builds by 15-30% or more
- docs.rs/hakari

Questions?